

# PACMAN + DQNs

Elijah Tai





# Contents

- Topic + Background
- Hypothesis + Importance
- Progress from Midterm Report
- Results + Code
- Issues + Solutions
- Next Steps



# Contents



Topic + Background



Hypothesis + Importance



Progress from Midterm Report



Results + Code



Issues + Solutions




Next Steps




$$Q^{new}(s_t, a_t) \leftarrow \underbrace{Q(s_t, a_t)}_{\text{current value}} + \underbrace{\alpha}_{\text{learning rate}} \cdot \underbrace{\left( \underbrace{r_t}_{\text{reward}} + \underbrace{\gamma}_{\text{discount factor}} \cdot \underbrace{\max_a Q(s_{t+1}, a)}_{\text{estimate of optimal future value}} - \underbrace{Q(s_t, a_t)}_{\text{current value}} \right)}_{\text{temporal difference}}$$

new value (temporal difference target)




$$Q^{new}(s_t, a_t) \leftarrow \underbrace{Q(s_t, a_t)}_{\text{current value}} + \underbrace{\alpha}_{\text{learning rate}} \cdot \underbrace{\left( \underbrace{r_t}_{\text{reward}} + \underbrace{\gamma}_{\text{discount factor}} \cdot \underbrace{\max_a Q(s_{t+1}, a)}_{\text{estimate of optimal future value}} - \underbrace{Q(s_t, a_t)}_{\text{current value}} \right)}_{\text{temporal difference}}$$

new value (temporal difference target)




$$Q_{t+1}^A(s_t, a_t) = Q_t^A(s_t, a_t) + \alpha_t(s_t, a_t) \left( r_t + \gamma Q_t^B \left( s_{t+1}, \arg \max_a Q_t^A(s_{t+1}, a) \right) - Q_t^A(s_t, a_t) \right)$$

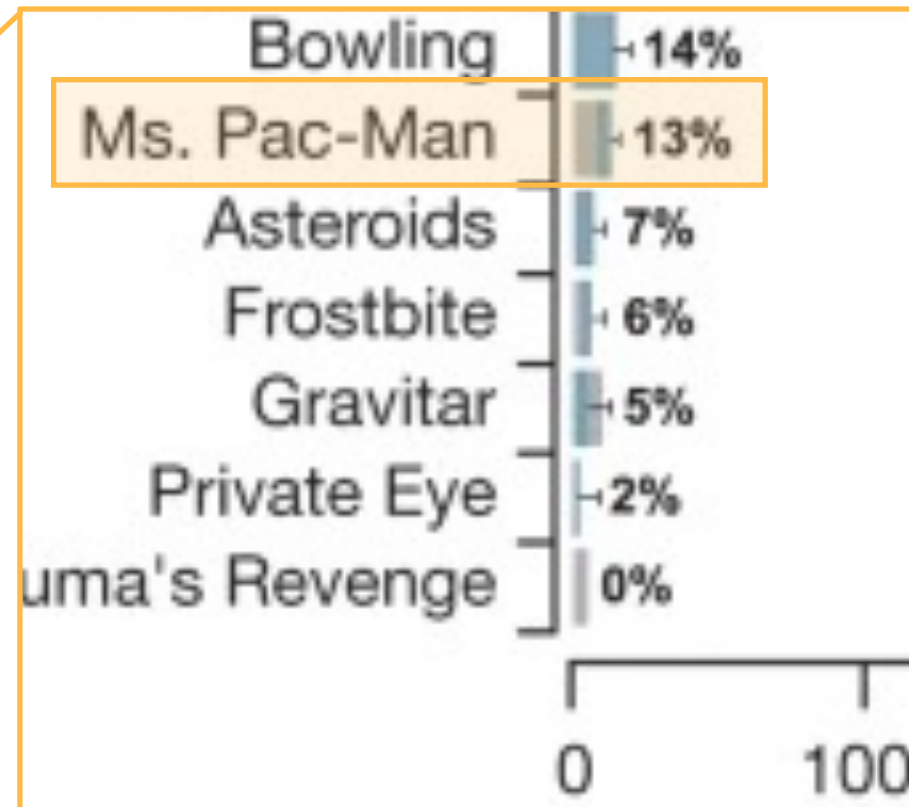



Watkins, C.J.C.H., Dayan, P. *Q*-learning. *Mach Learn* **8**, 279–292 (1992).  
<https://doi.org/10.1007/BF00992698>

Volodymyr Mnih et al. “Playing atari with deep reinforcement learning”.  
*arXiv preprint arXiv:1312.5602* (2013).



Mnih, V., Kavukcuoglu, K., Silver, D. *et al.* Human-level control through deep reinforcement learning. *Nature* **518**, 529–533 (2015).  
<https://doi.org/10.1038/nature14236>



Mnih, V., Kavukcuoglu, K., Silver, D. *et al.* Human-level control through deep reinforcement learning. *Nature* **518**, 529–533 (2015).  
<https://doi.org/10.1038/nature14236>

# Contents

Topic + Background



**Hypothesis + Importance**



Progress from Midterm Report



Results + Code



Issues + Solutions



Next Steps





# Hypothesis

- DQN algorithm struggles in MS PACMAN because of the power-up pellets. Suddenly, the best strategy is to chase ghosts instead of fleeing. This change confuses the network.

Removing the reward from eating ghosts during training paradoxically increases rewards.

Initializing powered-up MS PACMAN to do the opposite movement increases rewards.

Discretely context switching between two different networks will increase rewards.

Valid combinations of the other optimizations will further increase rewards.

# Importance

- While the optimizations are specific to MS PACMAN, this generally demonstrates tactics to improve performance, especially within settings of low computational resources when the neural network may not have a chance to stably converge.

# Contents

Topic + Background

Hypothesis + Importance

 **Progress from Midterm Report**

 Results + Code

 Issues + Solutions

 Next Steps



# Progress

- Previously, I got an existing DQN running from: [github.com/bourbonut/dqn-pacman](https://github.com/bourbonut/dqn-pacman)
- Learned about DQNs and fixed bugs in the original code.
- Wrote a script to determine the invincibility frames of ms pacman after eating a power-pellet
- Coded different behaviors triggered by certain score rewards are triggers.
- Ran 10k frames ~65 epochs on three of the four specific hypotheses
- Generated plots for visualization



# Contents

Topic + Background

Hypothesis + Importance

Progress from Midterm Report



**Results + Code**



Issues + Solutions

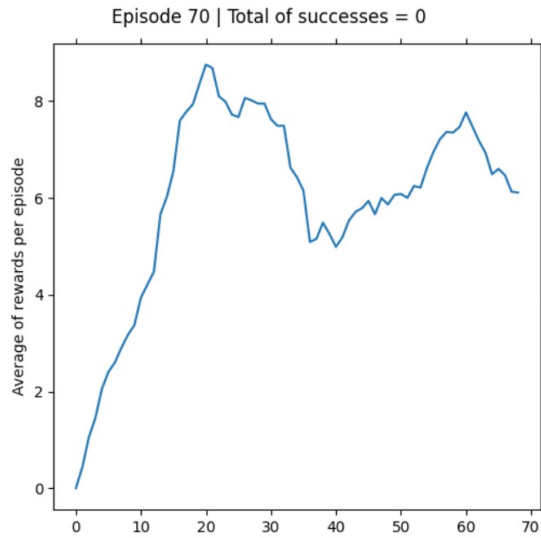


Next steps



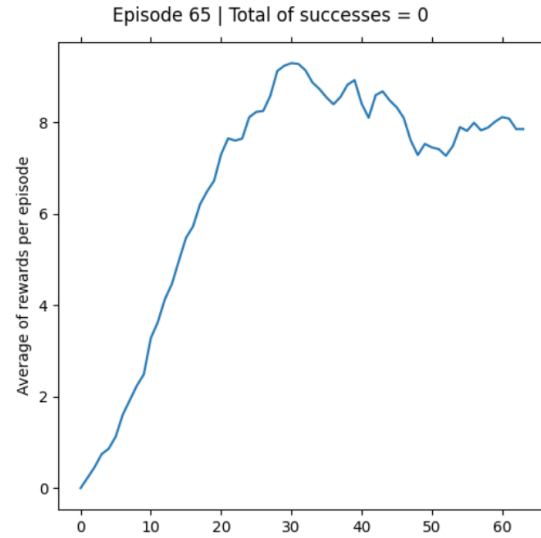
# Results: 10k steps

Neither



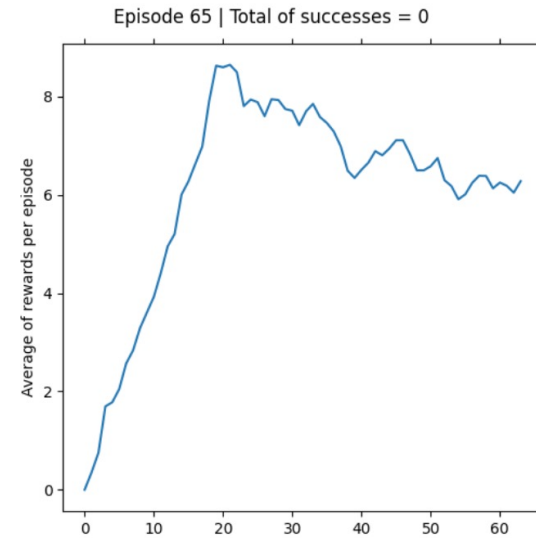
6.0

No Ghost  
Reward




7.8

Opposite  
Initialization



6.5


Both



8.1


# Results: 10k steps

Neither




0.054

No Ghost  
Reward




0.027

Opposite  
Initialization



0.040


Both



0.042


# Results: 10k steps

Neither




1.72

No Ghost  
Reward




1.47

Opposite  
Initialization



1.62

Both



1.70

# Main Code Alterations

## Invincibility Framing

```
1 import gym
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import torch
5 import random as r
6
7 env = gym.make("MsPacman-v4", render_mode='human')
8 env.reset()
9
10 stepmode = False
11 phase = 0
12 counter = 0
13 for _ in range(10000):
14     # env.action_space.n == 9
15     # env.get_action_meanings() == ['NOOP', 'UP', 'RIGHT', 'LEFT', 'DOWN']
16
17     if phase == 0:
18         next_state, reward, terminated, truncated, info = env.step(4)
19         counter += 1
20         if counter == 120:
21             phase += 1
22             counter = 0
23     if phase == 1:
24         next_state, reward, terminated, truncated, info = env.step(3)
25         counter += 1
26         if counter >= 80:
27             phase += 1
28             counter = 0
29     if phase == 2:
30         counter += 1
31         next_state, reward, terminated, truncated, info = env.step(1)
32         if reward == 50:
33             phase += 1
34             counter = 0
35     if phase == 3:
36         counter += 1
37         input(f'continue {counter}:')
38         next_state, reward, terminated, truncated, info = env.step(0)
39
40     env.render()
41     print(phase)
42
43 env.close()
```

## No Ghost Reward

```
16 parser.add_argument(
17     "--ghostbonus",
18     action="store_true",
19     dest="ghostbonus",
20     help="no reward for eating a ghost or a strawberry",
21 )
22
23
24 # Elijah: Change the reward to ignore eating a ghost
25 display_reward = reward_
26 if args.ghostbonus and (reward_ == 200 or reward_ % 400 == 0):
27     display_reward = 0
28
29
30 # Elijah: We want to display the true reward, but we want to
31 if args.ghostbonus:
32     display.data.rewards.append(display_reward)
33 else:
34     display.data.rewards.append(reward)
35 reward = torch.tensor([reward], device=device)
36
37 old_action = action_
38 if reward != 0:
39     dmaker.old_action = action.item()
40
41 next_state = preprocess_observation(observations, obs)
```

```
119 def _save(self, image=False):
120     """Save data in 'pickle' file and an image"""
121     if image:
122         PATH_PLOTS = self.data.path / '..' / 'plots'
123         self.update_axis()
124         self.fig.tight_layout()
125         plt.savefig(PATH_PLOTS / f"episode-{self.data.ep}.png")
126         print(f"Figure {self.data.ep} saved.")
127         for axis in self.axis:
128             axis.cla()
129         self.data.save()
```

## Opposite Initialization

```
22 parser.add_argument(
23     "--ghostrev",
24     action="store_true",
25     dest="ghostrev",
26     help="train a single NN, reversing actions under power pellet"
27 )
28
29
30 # Elijah: Reverse the action that we do
31 if args.ghostrev and rev_counter > 0:
32     action = [0, 4, 3, 2, 1, 9, 7, 6, 5][action]
33     rev_counter -= 1
34
35 action_ = ACTIONS[old_action][action.item()] # Elijah: This sort of softens the new action.
36 obs, reward_, done, info = env.step(action_)
```

```
20 REVERSED = (0, 1, 1, 0, 2, 3, 3, 2)
21 isreversed = 1
22 lambda last_action, action: "default" if REVERSED[action] == last_action else "reverse"
23
24
25 ACTIONS = {
26     1: [1, 4, 6, 5],
27     2: [5, 7, 3, 2],
28     3: [6, 8, 3, 2],
29     4: [1, 4, 6, 7],
30     5: [1, 4, 3, 2],
31     6: [1, 4, 3, 2],
32     7: [1, 4, 3, 2],
33     8: [1, 4, 3, 2],
34 }
```

```
22 TARGET_UPDATE = 400 # here, Elijah changed from 8
23 REPLAY_MEMORY_SIZE = 3 * 1200 #here, Elijah changed
24
25 # Environment constants
26 N_ACTIONS = 4
27 AVOIDED_STEPS = 80 # At the beginning, there is a
28 DEAD_STEPS = 36 # frames to avoid when the agent d
29 K_FRAME = 2
30
31 # Optimizer parameters
32 LEARNING_RATE = 2.5e-4
33 # DECAY_RATE = 0.99
34 MOMENTUM = 0.95
35
36 # Algorithm constant
37 MAX_FRAMES = 20_000 #Elijah changed from 2,000,000
38 SAVE_MODEL = 5 # Elijah changed from 20
```

```
ghostbonusDQN
├── models
│   ├── policy-model-5.pt
│   ├── policy-model-10.pt
│   ├── policy-model-15.pt
│   ├── policy-model-20.pt
│   ├── policy-model-25.pt
│   ├── policy-model-30.pt
│   ├── policy-model-35.pt
│   ├── policy-model-40.pt
│   ├── policy-model-45.pt
│   ├── policy-model-50.pt
│   ├── policy-model-55.pt
│   ├── policy-model-60.pt
│   ├── policy-model-65.pt
│   └── policy-model-final.pt
├── target-model-5.pt
├── target-model-10.pt
├── target-model-15.pt
├── target-model-20.pt
├── target-model-25.pt
├── target-model-30.pt
├── target-model-35.pt
├── target-model-40.pt
├── target-model-45.pt
├── target-model-50.pt
├── target-model-55.pt
├── target-model-60.pt
├── target-model-65.pt
└── target-model-final.pt
```

```
plots
├── episode-5.png
├── episode-10.png
├── episode-15.png
├── episode-20.png
├── episode-25.png
├── episode-30.png
├── episode-35.png
├── episode-40.png
├── episode-45.png
├── episode-50.png
├── episode-55.png
├── episode-60.png
├── episode-65.png
├── episode-5.pkl
├── episode-10.pkl
├── episode-15.pkl
├── episode-20.pkl
├── episode-25.pkl
├── episode-30.pkl
├── episode-35.pkl
├── episode-40.pkl
├── episode-45.pkl
├── episode-50.pkl
├── episode-55.pkl
├── episode-60.pkl
└── episode-65.pkl
```



# Demo



# Contents

Topic + Background

Hypothesis + Importance

Progress from Midterm Report

Results + Code

**Issues + Solutions**

Next Steps



# Issues + Solutions

- Need variance metric for comparing the learning curves
  - Need to run the code multiple times to make better comparison
- Only have results for short time frames
  - Plan to let it run for  $\sim 1000$  episodes
- Still need to add in the context switching (double) double DQN
  - Requires a new class and code structure changes
- This will add more quantitative rigor



# Contents

Topic + Background

Hypothesis + Importance

Progress from Midterm Report

Results + Code

Issues + Solutions

Next Steps



# Planning for Next Steps

- Running more times for longer
- Implementing the Double-Double DQN (the third hypothesis)
- Generating more quantitative metrics for learning curve comparison
- Extra: Implementing Dueling DQN or soft steps (Polyak averaging)
- Extra: Comparing greyscale versus color to understand ghost personalities





Thank you!

